


```

# Define the optimizer (SGD with learning rate and momentum)
optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.5)
# Create criterion
print("Optimizer completed!")
print("Criterion: CrossEntropyLoss")
print("Optimizer: SGD (lr=0.01, momentum=0.5)")

# Model Architecture:
model = nn.Sequential(
    # (x): Linear(in_features=768, out_features=128, bias=True)
    # (x): Linear(in_features=128, out_features=64, bias=True)
    # (x): Linear(in_features=64, out_features=32, bias=True)
)
Model has 49,386 trainable parameters
Sgd completed.
Criterion: CrossEntropyLoss
Optimizer: SGD (lr=0.01, momentum=0.5)

Finally you can define a training, and test loop

# Create an array to log the loss and accuracy
train_losses = []
train_steps = []
test_steps = []
test_accuracy = []
current_step = 0 # Start with global step 0
current_epoch = 0 # Start with epoch 0
# Define the number of epochs
epochs = 10
# Train and testing loop
for epoch in range(epochs):
    print(f"Epoch {epoch+1}/{epochs}")
    print(f"Training Loss: {(epoch_train_loss / len(train_loader)), 4f}")
    # Training Phase
    model.train() # Set the model to training mode
    epoch_train_loss = 0
    for batch_idx, (data, target) in enumerate(train_loader):
        optimizer.zero_grad() # Zero the gradients
        output = model(data) # Forward pass
        criterion(output, target) # Calculate loss
        loss.backward() # Backward pass
        optimizer.step() # Optimization step
        epoch_train_loss += loss.item()
        epoch_train_steps.append(current_step)
        current_step += 1
    print(f"Training loss: {(epoch_train_loss / len(train_loader)), 4f}")
    # Testing Phase
    model.eval() # Set the model to evaluation mode
    epoch_eval_loss = 0
    correct_predictions = 0
    total_samples = 0
    with torch.no_grad(): # Disable gradient computation
        for data, target in test_loader:
            output = model(data) # Forward pass
            loss = criterion(output, target) # Calculate loss
            epoch_eval_loss += loss.item()
            pred = output.argmax(dim=1) # Get the class index with the highest score
            correct_predictions += (pred == target).sum().item() # Count correct predictions
            total_samples += target.size(0) # Count total samples
    epoch_test_steps.append(epoch * len(test_loader))
    test_losses.append(epoch_eval_loss / len(test_loader))
    test_steps.append(total_samples / len(test_loader))
    test_accuracy.append(correct_predictions / total_samples)
    print(f"Test Loss: {(epoch_eval_loss / len(test_loader)), 4f}")
    print(f"Test Accuracy: {(correct_predictions / total_samples), 4f}%")


# Epoch 1/10
# Training Loss: 0.5587
# Test Loss: 0.2763
# Test Accuracy: 0.9200
# Epoch 2/10
# Training Loss: 0.2482
# Test Loss: 0.2095
# Test Accuracy: 0.9399
# Epoch 3/10
# Training Loss: 0.1392
# Test Loss: 0.1034
# Test Accuracy: 0.9513
# Epoch 4/10
# Training Loss: 0.1488
# Test Loss: 0.1484
# Test Accuracy: 0.9578
# Epoch 5/10
# Training Loss: 0.1226
# Test Loss: 0.0930
# Test Accuracy: 0.9637

# Declare the train function
def cpu_train(epoch, train_losses, steps, current_step):
    model.train()
    ph = enumerate(train_loader, total=len(train_loader))
    for batch_idx, (data, target) in ph:
        optimizer.zero_grad()
        output = model(data)
        loss = criterion(output, target)
        loss.backward()
        optimizer.step()
        current_step += 1
        if batch_idx % 10 == 0:
            steps.append(current_step)
            desc = f'Train (epoch: {epoch}) [{batch_idx * len(data)}:{len(train_loader.dataset)}] loss: {loss.item():.4f}'
            ph.set_description(desc)
    return current_step
# Define the test function
def cpu_test(epoch, test_losses, test_accuracy, steps, current_step):
    model.eval()
    test_steps.append(steps[-1])
    correct = 0
    ph = enumerate(test_loader, total=len(test_loader), desc="Testing...")
    with torch.no_grad():
        for data, target in ph:
            output = model(data)
            pred = output.argmax(dim=1, keepdim=True)
            correct += pred.eq(target.view_as(pred)).sum().item()
    test_losses.append(len(test_loader) * epoch * len(data) / len(test_loader))
    test_accuracy.append(correct / len(test_loader.dataset))
    steps.append(current_step)
    print(f"({epoch + 1}, {correct / len(test_loader.dataset)}, {test_accuracy[-1]})")
    # Train for 10 epochs
    current_epoch = 0
    for epoch in range(10):
        current_epoch = cpu_train(epoch, train_losses, train_steps, current_epoch)
        print(f"Current epoch: {epoch+1}/{10}, Loss: {train_losses[-1]}, Accuracy: {train_accuracy[-1]}%")
    print(f"Final epoch: {current_epoch}, Loss: {train_losses[-1]}, Accuracy: {train_accuracy[-1]}%")

Train Epoch: 0 [57600/50000 (98%) Loss: 0.4052±0.0001] 938/938 [00:12:00±00, 75.14it/s]
Testing... 3M0K [17/157 [00:02:00±00, 70.72it/s]
Test set: Average loss: 0.0984, Accuracy: 96.21/10000 (97)
Train Epoch: 1 [57600/50000 (98%) Loss: 0.0384±0.0001] 938/938 [00:12:00±00, 72.42it/s]
Testing... 3M0K [17/157 [00:02:00±00, 75.22it/s]
Test set: Average loss: 0.0091, Accuracy: 98.92/10000 (97)
Train Epoch: 2 [57600/50000 (98%) Loss: 0.0292±0.0001] 938/938 [00:12:00±00, 73.95it/s]
Testing... 3M0K [17/157 [00:02:00±00, 80.38it/s]
Test set: Average loss: 0.0015, Accuracy: 97.24/10000 (97)
Train Epoch: 3 [57600/50000 (98%) Loss: 0.0179±0.0001] 938/938 [00:12:00±00, 74.14it/s]
Testing... 3M0K [17/157 [00:02:00±00, 80.24it/s]
Test set: Average loss: 0.0000, Accuracy: 97.98/10000 (97)
Train Epoch: 4 [57600/50000 (98%) Loss: 0.0153±0.0001] 938/938 [00:13:00±00, 70.72it/s]
Testing... 3M0K [17/157 [00:03:00±00, 81.57it/s]
Test set: Average loss: 0.0050, Accuracy: 97.42/10000 (97)
Train Epoch: 5 [57600/50000 (98%) Loss: 0.0052±0.0001] 938/938 [00:12:00±00, 75.76it/s]
Testing... 3M0K [17/157 [00:02:00±00, 77.89it/s]
Test set: Average loss: 0.0037, Accuracy: 97.22/10000 (97)
Train Epoch: 6 [57600/50000 (98%) Loss: 0.0045±0.0001] 938/938 [00:12:00±00, 74.46it/s]
Testing... 3M0K [17/157 [00:02:00±00, 67.66it/s]
Test set: Average loss: 0.0013, Accuracy: 97.52/10000 (97)
Train Epoch: 7 [57600/50000 (98%) Loss: 0.0044±0.0001] 938/938 [00:12:00±00, 75.35it/s]
Testing... 3M0K [17/157 [00:03:00±00, 60.53it/s]
Test set: Average loss: 0.0074, Accuracy: 97.68/10000 (97)
Train Epoch: 8 [57600/50000 (98%) Loss: 0.0040±0.0001] 938/938 [00:12:00±00, 74.99it/s]
Testing... 3M0K [17/157 [00:02:00±00, 60.64it/s]
Test set: Average loss: 0.0037, Accuracy: 97.76/10000 (97)
Train Epoch: 9 [57600/50000 (98%) Loss: 0.0038±0.0001] 938/938 [00:12:00±00, 74.12it/s]
Testing... 3M0K [17/157 [00:02:00±00, 65.36it/s]
Test set: Average loss: 0.0015, Accuracy: 97.73/10000 (97)
Train Epoch: 10 [57600/50000 (98%) Loss: 0.0035±0.0001] 938/938 [00:12:00±00, 75.01it/s]
Testing... 3M0K [17/157 [00:03:00±00, 64.63it/s]
Test set: Average loss: 0.0027, Accuracy: 97.72/10000 (97)
Train Epoch: 11 [57600/50000 (98%) Loss: 0.0034±0.0001] 938/938 [00:12:00±00, 74.35it/s]
Testing... 3M0K [17/157 [00:03:00±00, 60.32it/s]
Test set: Average loss: 0.0013, Accuracy: 97.80/10000 (97)

# Train for 10 epochs
for epoch in range(10):
    current_step = cpu_train(current_epoch, train_losses, train_steps, current_step)
    print(f"Current epoch: {epoch+1}/{10}, Loss: {train_losses[-1]}, Accuracy: {train_accuracy[-1]}%")
    current_epoch += 1

Train Epoch: 0 [57600/50000 (98%) Loss: 0.0320±0.0001] 938/938 [00:12:00±00, 75.22it/s]
Testing... 3M0K [17/157 [00:02:00±00, 85.36it/s]
Test set: Average loss: 0.0723, Accuracy: 97.06/10000 (98)
Train Epoch: 0 [57600/50000 (98%) Loss: 0.0316±0.0001] 938/938 [00:12:00±00, 74.12it/s]
Testing... 3M0K [17/157 [00:02:00±00, 85.73it/s]
Test set: Average loss: 0.0715, Accuracy: 97.23/10000 (98)
Train Epoch: 0 [57600/50000 (98%) Loss: 0.0325±0.0001] 938/938 [00:12:00±00, 75.01it/s]
Testing... 3M0K [17/157 [00:03:00±00, 84.63it/s]
Test set: Average loss: 0.0723, Accuracy: 97.22/10000 (98)
Train Epoch: 0 [57600/50000 (98%) Loss: 0.0315±0.0001] 938/938 [00:12:00±00, 74.35it/s]
Testing... 3M0K [17/157 [00:03:00±00, 80.32it/s]
Test set: Average loss: 0.0713, Accuracy: 97.80/10000 (97)

```

Train Epoch: 3984 [57000/68800 (98%)] Loss: 0.083407, 1000ms
Testing...: 1000X 157/157 [0.08:01:00] 65,291/65,291 (100%)
Test set: Average loss: 0.075, Accuracy: 97/51 (100%)

Train Epoch: 3985 [57000/68800 (98%)] Loss: 0.080986, 1000ms
Testing...: 1000X 157/157 [0.08:01:00] 86,351/86,351 (100%)
Test set: Average loss: 0.075, Accuracy: 97/78 (100%)

Train Epoch: 3986 [57000/68800 (98%)] Loss: 0.081228, 1000ms
Testing...: 1000X 157/157 [0.08:01:00] 86,351/86,351 (100%)
Test set: Average loss: 0.075, Accuracy: 97/78 (100%)

Train Epoch: 3987 [57000/68800 (98%)] Loss: 0.078337, 1000ms
Testing...: 1000X 157/157 [0.08:01:00] 86,351/86,351 (100%)
Test set: Average loss: 0.075, Accuracy: 97/80 (100%)

Train Epoch: 3988 [57000/68800 (98%)] Loss: 0.072598, 1000ms
Testing...: 1000X 157/157 [0.08:01:00] 86,351/86,351 (100%)
Test set: Average loss: 0.078, Accuracy: 97/80 (100%)

Train Epoch: 3989 [57000/68800 (98%)] Loss: 0.071089, 1000ms
Testing...: 1000X 157/157 [0.08:01:00] 86,351/86,351 (100%)
Test set: Average loss: 0.075, Accuracy: 97/77 (100%)

Question 2

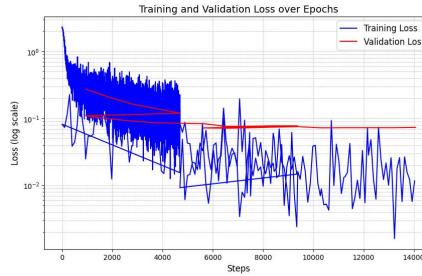
Using the skills you acquired in the previous assignment edit the cell below to use matplotlib to visualize the loss for training and validation for the first 10 epochs. They should be plotted on the same graph, labeled, and use a log-scale on the y-axis.

```

# visualize training and validation loss
# Visualize training and validation loss
def plot_losses(train_losses, train_steps, test_losses, test_steps):
    plt.figure(figsize=(10, 6))
    # Plot training loss
    plt.plot(train_steps, train_losses, label="Training loss", color='blue', linewidth=1.5)
    # Plot validation loss
    plt.plot(test_steps, test_losses, label="Validation loss", color='red', linewidth=1.5)
    # Set the x-axis to log scale
    plt.xscale('log')
    # Add labels, title, and legend
    plt.title("Training and Validation loss over Epochs", fontsize=14)
    plt.xlabel("Epochs", fontsize=12)
    plt.ylabel("Loss (log Scale)", fontsize=12)
    plt.legend(['Training and Validation loss over Epochs', fontsize=14])
    plt.grid(True, which="both", linestyle="--", linewidth=0.5)
    plt.show()

# Call the function to visualize the data
plot_losses(train_losses, train_steps, test_losses, test_steps)

```



Question 3

The model may be able to train for a bit longer. Edit the cell below to modify the previous training code to also report the time per epoch and the time for 10 epochs with testing. You can use `time.time()` to get the current time in seconds. Then run the model for another 10 epochs, printing out the execution time at the end, and replot the loss functions with the extra 10 epochs below.

```

# initialize the timer
t = time.time()
start_t = t

# Initialize the timer
t = time.time()

# Train for 10 additional epochs and track time per epoch
for epoch in range(ep_start, ep_end):
    # Compute the time taken for the previous epoch
    current_time = time.time()
    epoch_time = current_time - start_t
    start_t = current_time

    # Train and test for the current epoch
    train(ep_start, epoch)
    current_time = time.time()
    train_losses, train_accs, train_tstep, current_tstep = get_train_info(ep_start, epoch)

    # Print training accuracy, test accuracy, time taken, current epoch
    print("Epoch %d: Accuracy: %f, Accuracy: %f, Time: %f seconds" % (epoch, current_accuracy, test_accuracy, epoch_time))

    # Increment the epoch counter
    current_epoch += 1

# Print the total time for this epoch
epoch_time = time.time() - epoch_start
print("Epoch %d: Time: %f seconds" % (ep_end, epoch_time))

# Print the total time for all epochs
total_time = time.time() - start_t
print("Total time for %d epochs: (%f, %f) seconds" % (ep_end, total_time, total_time))

```

```
Train Epoch: 0 [57000/58000 (98%)] Loss: 0.086731: 1000 [938/938 [08:12:00, 80, 74.961it/T]
Testing...: 1000 [157/157 [08:12:00, 80, 76.817it/T]
Test set: Average loss: 0.0792, Accuracy: 9777/10000 (98%)
Epoch 11 completed in 14.48 seconds
Train Epoch: 11 [57000/58000 (98%)] Loss: 0.166651: 1000 [938/938 [08:12:00, 80, 74.601it/T]
Testing...: 1000 [157/157 [08:12:00, 80, 76.814it/T]
Test set: Average loss: 0.0774, Accuracy: 9777/10000 (98%)
Epoch 12 completed in 14.88 seconds
Train Epoch: 12 [57000/58000 (98%)] Loss: 0.166651: 1000 [938/938 [08:12:00, 80, 75.471it/T]
Testing...: 1000 [157/157 [08:12:00, 80, 76.814it/T]
Test set: Average loss: 0.0731, Accuracy: 9780/10000 (98%)
Epoch 13 completed in 14.92 seconds
Train Epoch: 13 [57000/58000 (98%)] Loss: 0.161556: 1000 [938/938 [08:13:00, 80, 79.221it/T]
Testing...: 1000 [157/157 [08:13:00, 80, 76.351it/T]
Test set: Average loss: 0.0753, Accuracy: 9780/10000 (98%)
Epoch 14 completed in 14.98 seconds
Train Epoch: 14 [57000/58000 (98%)] Loss: 0.081066: 1000 [938/938 [08:13:00, 80, 74.831it/T]
Testing...: 1000 [157/157 [08:13:00, 80, 76.351it/T]
Test set: Average loss: 0.0752, Accuracy: 9780/10000 (98%)
Epoch 15 completed in 14.19 seconds
Train Epoch: 15 [57000/58000 (98%)] Loss: 0.084959: 1000 [938/938 [08:12:00, 80, 74.461it/T]
Testing...: 1000 [157/157 [08:12:00, 80, 76.351it/T]
Test set: Average loss: 0.0752, Accuracy: 9780/10000 (98%)
Epoch 16 completed in 14.41 seconds
Train Epoch: 16 [57000/58000 (98%)] Loss: 0.080173: 1000 [938/938 [08:12:00, 80, 75.481it/T]
Testing...: 1000 [157/157 [08:12:00, 80, 65.961it/T]
Test set: Average loss: 0.0772, Accuracy: 9780/10000 (98%)
Epoch 17 completed in 14.19 seconds
Train Epoch: 17 [57000/58000 (98%)] Loss: 0.086071: 1000 [938/938 [08:12:00, 80, 75.841it/T]
Testing...: 1000 [157/157 [08:12:00, 80, 67.461it/T]
Test set: Average loss: 0.0772, Accuracy: 9783/10000 (98%)
Epoch 18 completed in 14.31 seconds
Train Epoch: 18 [57000/58000 (98%)] Loss: 0.087933: 1000 [938/938 [08:12:00, 80, 75.131it/T]
Testing...: 1000 [157/157 [08:12:00, 80, 76.961it/T]
Test set: Average loss: 0.0860, Accuracy: 9779/10000 (98%)
Epoch 19 completed in 14.15 seconds
Train Epoch: 19 [57000/58000 (98%)] Loss: 0.081771: 1000 [938/938 [08:12:00, 80, 74.821it/T]
Testing...: 1000 [157/157 [08:12:00, 80, 84.841it/T]
Test set: Average loss: 0.0878, Accuracy: 9777/10000 (98%)
```

Question 4

Make an observation from the above plot. Do the test and train loss curves indicate that the model should train longer to improve accuracy? Does it indicate that 20 epochs is too long? Edit the cell below to answer these questions.

Double-click (or enter) to edit

▼ Moving to the GPU

Using a GPU with torch is relatively simple, but has a few gotchas. Torch abstracts away most of the CUDA n

Using a GPU with torch is relatively simple, but has a few gotchas. Torch abstracts away most of the CUDA runtime API, but it has a few hold-over concepts such as moving data between devices. Additionally, since the GPU is treated as a device separate from the CPU, you cannot combine CPU and GPU based tensors in the same operation. Doing so will result in a device mismatch error. If this occurs, check where the tensors are located (you can always print ..device on a tensor), and make sure they have been properly moved to the correct device.

You will start by creating a new model, optimizer and criterion (not really necessary in this case since you already did this above but it's better).

you will start by creating a model, optimising, and evaluating (not really necessary in this case, since you'll be using Colab's GPU and the device is set to "GPU" by default). However, for clarity and completeness), however, one change that you'll make is moving the model to the GPU first. This can be done by calling `.cuda()` in general, or `.to("cuda")` to be more explicit. In general specific GPU devices can be targeted such as `.to("cuda:0")` for the first GPU (index 0), etc., but since there is only one GPU in Colab this is not necessary in this case.

```

# create the model
model = MLP()

# move the model to the GPU
model.cuda()

# for a criterion (loss) function, we will use Cross-Entropy loss. This is the most common criterion used for multi-class prediction, and is also used by tokenized transformer models

```

```
# Which is < 0. For MNIST, the prediction right be [-0.0056, -0.2044, 1.1736, 0.0859, 1.8443, -0.9627, 0.0785, -1.0752, 1.1376, 1.0320], with the label 3.
# Cross-entropy can be thought of as finding the difference between what the predicted distribution and the one-hot distribution

criterion = nn.CrossEntropyLoss()

# Then you can instantiate the optimizer. You will use Stochastic Gradient Descent (SGD), and can set the learning rate to 0.1 with a momentum factor of 0.5
# The first input to the optimizer is the list of model parameters, which is obtained by calling .parameters() on the model object
optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.5)

import torch
import torch.nn as nn
import torch.optim as optim
import torch.onnx._functional as F

# Create a new array to log the loss and accuracy
train_losses = []
train_steps = []
test_steps = []
test_losses = []
test_steps = []
test_accuracy = []

current_step = 0 # Start with global step 0
current_epoch = 0 # Start with epoch 0

# Set up the device for GPU or CPU
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Create the model and move it to the device
model = MLP().to(device)

# Create the criterion (loss function)
criterion = nn.CrossEntropyLoss()

# Create the optimizer
optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.5)

# Training and testing loop
for epoch in range(10): # For 10 epochs (adjust as necessary)
    model.train() # Set the model to training mode
    current_epoch += 1
    correct_train = 0
    total_train = 0

    for i, (inputs, labels) in enumerate(train_loader): # Assuming train_loader exists
        inputs, labels = inputs.to(device), labels.to(device) # Move data to device
        optimizer.zero_grad() # Zero gradients for each batch

        outputs = model(inputs) # Forward pass
        loss = criterion(outputs, labels) # Compute loss
        loss.backward() # Backpropagates the loss
        optimizer.step() # Updates weights

        # Track loss and accuracy
        running_loss += loss.item()
        total_train += labels.size(0)
        correct_train += (predicted == labels).sum().item()

        # Log every 100 steps (adjust as needed)
        if i % 100 == 0:
            print(f"Step [{i}/{len(train_loader)}], Loss: {loss.item():.4f}")

    # Average loss and accuracy for the epoch
    train_losses.append(running_loss / len(train_loader))
    train_accuracy = 100 * correct_train / total_train
    train_steps.append(current_step)

    print(f"Epoch [{epoch+1}/10], Train Loss: {train_losses[-1]:.4f}, Train Accuracy: {train_accuracy:.2f}%")

    # Testing loop
    model.eval() # Set the model to evaluation mode
    running_test_loss = 0.0
    correct_test = 0
    total_test = 0

    with torch.no_grad(): # No gradient calculation during testing
        for inputs, labels in test_loader: # Assuming test_loader exists
            inputs, labels = inputs.to(device), labels.to(device) # Move data to device

            outputs = model(inputs)
            loss = criterion(outputs, labels)
            running_test_loss += loss.item()

            predicted = torch.argmax(outputs.data, 1)
            total_test += labels.size(0)
            correct_test += (predicted == labels).sum().item()

    # Average test loss and accuracy
    test_losses.append(running_test_loss / len(test_loader))
    test_accuracy_value = 100 * correct_test / total_test
    test_steps.append(current_step)

    print(f"Test Loss: {test_losses[-1]:.4f}, Test Accuracy: {test_accuracy_value:.2f}%")

    current_epoch += 1 # Increment the epoch count
    current_step += 1 # Increment the global step
```

```
Step 0/938, Loss: 2.0000
Step 180/938, Loss: 1.2485
Step 280/938, Loss: 0.4886
Step 380/938, Loss: 0.3980
Step 480/938, Loss: 0.3980
Step 580/938, Loss: 0.3980
Step 680/938, Loss: 0.2888
Step 780/938, Loss: 0.1372
Step 880/938, Loss: 0.1360
Step 980/938, Loss: 0.3460
Epoch 1/10, Train Loss: 0.2458, Train Accuracy: 86.72%
Train Loss: 0.2458, Train Accuracy: 86.72%
Step 0/938, Loss: 0.1797
Step 180/938, Loss: 0.1700
Step 280/938, Loss: 0.3102
Step 380/938, Loss: 0.3855
Step 480/938, Loss: 0.1688
Step 580/938, Loss: 0.1688
Step 680/938, Loss: 0.1688
Step 780/938, Loss: 0.1630
Step 880/938, Loss: 0.2130
Step 980/938, Loss: 0.1630
Epoch 2/10, Train Loss: 0.2458, Train Accuracy: 92.88%
Train Loss: 0.2458, Train Accuracy: 92.88%
Test Loss: 0.1867, Test Accuracy: 94.62%
Step 0/938, Loss: 0.1673
Step 180/938, Loss: 0.1583
Step 280/938, Loss: 0.1583
Step 380/938, Loss: 0.1583
Step 480/938, Loss: 0.1583
Step 580/938, Loss: 0.1583
Step 680/938, Loss: 0.1583
Step 780/938, Loss: 0.1583
Step 880/938, Loss: 0.1583
Step 980/938, Loss: 0.1583
Epoch 3/10, Train Loss: 0.1546, Train Accuracy: 95.61%
Train Loss: 0.1546, Train Accuracy: 95.61%
Test Loss: 0.1673, Test Accuracy: 95.00%
Step 0/938, Loss: 0.1588
Step 180/938, Loss: 0.0943
Step 280/938, Loss: 0.1588
Step 380/938, Loss: 0.1588
Step 480/938, Loss: 0.1588
Step 580/938, Loss: 0.1588
Step 680/938, Loss: 0.1588
Step 780/938, Loss: 0.1588
Step 880/938, Loss: 0.1588
Step 980/938, Loss: 0.1588
Epoch 4/10, Train Loss: 0.1546, Train Accuracy: 95.61%
Train Loss: 0.1546, Train Accuracy: 95.61%
Test Loss: 0.1673, Test Accuracy: 95.00%
Step 0/938, Loss: 0.1588
Step 180/938, Loss: 0.0943
Step 280/938, Loss: 0.1588
Step 380/938, Loss: 0.1588
Step 480/938, Loss: 0.1588
Step 580/938, Loss: 0.1588
Step 680/938, Loss: 0.1588
Step 780/938, Loss: 0.1588
Step 880/938, Loss: 0.1588
Step 980/938, Loss: 0.1588
```

Now, copy your previous training code with the timing parameters below. It needs to be slightly modified to move everything to the GPU.

Before the line `output = model(data)`, add:

```
data = data.cuda()
target = target.cuda()
```

Note that this is needed in both the train and test functions.

Question 5

Please edit the cell below to show the new GPU train and test functions.

```
import time
import torch

# GPU train function
def gpu_train(epoch, train_losses, train_steps, current_step):
    model.train() # Set the model to training mode
    running_loss = 0.0
    correct_train = 0
    total_train = 0

    start_time_epoch = time.time() # Start time for this epoch

    for i, (data, target) in enumerate(train_loader): # Assuming train_loader exists
        data, target = data.to(device), target.to(device) # Move data to GPU

        optimizer.zero_grad() # Zero gradients for each batch
        output = model(data) # A forward pass
        loss = criterion(output, target) # Compute loss
        loss.backward() # Compute the loss
        optimizer.step() # Update the weights

        # Track loss and accuracy
        running_loss += loss.item()
        total_train += target.size(0)
        correct_train += (predicted == target).sum().item()

        # Log every 100 steps (adjust as needed)
        if i % 100 == 0:
            print(f"Step [{i}/{len(train_loader)}], Loss: {loss.item():.4f}")

    # Average loss and accuracy for the epoch
    train_losses.append(running_loss / len(train_loader))
    train_accuracy = 100 * correct_train / total_train
    train_steps.append(current_step)

    print(f"Epoch [{epoch+1}/{10}], Train Loss: {train_losses[-1]:.4f}, Train Accuracy: {train_accuracy:.2f}%")
```

```
# Time for the current epoch
current_time_epoch = time.time()
epoch_start_time = start_time_epoch
print(f"Epoch {epoch} completed in {(epoch_duration):.2f} seconds")

current_step += 1 # Increment the global step
return current_step

# GPU test function
def gpu_test(test_losses, test_accuracy, test_steps, current_step):
    model.eval() # Set the model to evaluation mode
    running_mean_loss = 0.0
    correct = 0
    total_test = 0

    with torch.no_grad(): # No gradient calculation during testing
        for data, target in test_loader: # Assuming test_loader exists
            data, target = data.to(device), target.to(device) # Move data to GPU
            output = model(data)
            loss = criterion(output, target)
            running_mean_loss += loss.item()

            _, predicted = torch.max(output.data, 1)
            total_test += target.size(0)
            correct += (predicted == target).sum().item()

    # Average test loss and accuracy
    test_losses.append(running_mean_loss / len(test_loader))
    test_accuracy_value = 100 * correct / total_test
    test_accuracy.append(test_accuracy_value)
    test_steps.append(current_step)

    print(f"Test Loss: {test_losses[-1]:.4f}, Test Accuracy: {test_accuracy_value:.2f}%")
    return test_accuracy_value

# Run training and testing for 20 epochs
current_step = 0
current_epoch = 0
train_losses = []
train_steps = []
test_steps = []
test_losses = []
test_accuracy = []

for epoch in range(20): # Modify number of epochs as required
    current_step, train_mean_loss, train_steps, current_step = gpu_train(test_losses, test_accuracy, test_steps, current_step)
    test_accuracy_value = gpu_test(test_losses, test_accuracy, test_steps, current_step)
    current_epoch += 1 # Increment the epoch count

    # Step [0/938], Loss: 0.0001
    # Step [100/938], Loss: 0.0015
    # Step [200/938], Loss: 0.0002
    # Step [300/938], Loss: 0.0001
    # Step [400/938], Loss: 0.0029
    # Step [500/938], Loss: 0.0020
    # Step [600/938], Loss: 0.0019
    # Step [700/938], Loss: 0.0579
    # Step [800/938], Loss: 0.0059
    # Step [900/938], Loss: 0.0050
    # Epoch [1/20], Train loss: 0.0054, Train Accuracy: 98.36%
    # Epoch 1 completed in 11.96 seconds
    # Test loss: 0.0018, Test Accuracy: 97.42%
    # Step [100/938], Loss: 0.0284
    # Step [200/938], Loss: 0.0133
    # Step [300/938], Loss: 0.0124
    # Step [400/938], Loss: 0.0242
    # Step [500/938], Loss: 0.0038
    # Step [600/938], Loss: 0.0037
    # Step [700/938], Loss: 0.0032
    # Step [800/938], Loss: 0.0032
    # Step [900/938], Loss: 0.0036
    # Epoch [2/20], Train loss: 0.0051, Train Accuracy: 98.33%
    # Epoch 2 completed in 11.96 seconds
    # Test loss: 0.0011, Test Accuracy: 97.64%
    # Step [100/938], Loss: 0.0530
    # Step [200/938], Loss: 0.0047
    # Step [300/938], Loss: 0.0116
    # Step [400/938], Loss: 0.0116
    # Step [500/938], Loss: 0.0157
    # Step [600/938], Loss: 0.0266
    # Step [700/938], Loss: 0.0060
    # Step [800/938], Loss: 0.0060
    # Step [900/938], Loss: 0.0118
    # Epoch [3/20], Train loss: 0.0045, Train Accuracy: 98.74%
    # Epoch 3 completed in 11.96 seconds
    # Test loss: 0.0009, Test Accuracy: 97.64%
    # Step [100/938], Loss: 0.0118
    # Step [200/938], Loss: 0.0019
    # Step [300/938], Loss: 0.0176
    # Step [400/938], Loss: 0.0061
    # Step [500/938], Loss: 0.0121
    # Step [600/938], Loss: 0.0321
    # Step [700/938], Loss: 0.0144
    # Step [800/938], Loss: 0.0144
    # Step [900/938], Loss: 0.0046
    # Epoch [4/20], Train loss: 0.0046, Train Accuracy: 98.85%
    # Epoch 4 completed in 11.96 seconds
    # Test loss: 0.0012, Test Accuracy: 97.64%
    # Step [100/938], Loss: 0.0363
    # Step [200/938], Loss: 0.1382
    # Step [300/938], Loss: 0.0568
    # Step [400/938], Loss: 0.0568
    # Step [500/938], Loss: 0.0154

    # One day training for 10 epochs
    # Run training and testing for 10 epochs
    current_step = 0
    current_epoch = 0
    train_losses = []
    train_steps = []
    test_steps = []
    test_losses = []
    test_accuracy = []

    for epoch in range(10): # Train for 10 epochs
        current_step, train_mean_loss, train_steps, current_step = gpu_train(test_losses, test_accuracy, test_steps, current_step)
        test_accuracy_value = gpu_test(test_losses, test_accuracy, test_steps, current_step)
        current_epoch += 1 # Increment the epoch count

        # Step [0/938], Loss: 0.0025
        # Step [100/938], Loss: 0.0075
        # Step [200/938], Loss: 0.0078
        # Step [300/938], Loss: 0.0078
        # Step [400/938], Loss: 0.0122
        # Step [500/938], Loss: 0.0167
        # Step [600/938], Loss: 0.0084
        # Step [700/938], Loss: 0.0121
        # Step [800/938], Loss: 0.0004
        # Step [900/938], Loss: 0.0004
        # Epoch [1/20], Train loss: 0.0008, Train Accuracy: 99.92%
        # Epoch 1 completed in 11.96 seconds
        # Test loss: 0.0001, Test Accuracy: 97.95%
        # Step [0/938], Loss: 0.0050
        # Step [100/938], Loss: 0.0000
        # Step [200/938], Loss: 0.0029
        # Step [300/938], Loss: 0.0057
        # Step [400/938], Loss: 0.0044
        # Step [500/938], Loss: 0.0013
        # Step [600/938], Loss: 0.0013
        # Step [700/938], Loss: 0.0013
        # Step [800/938], Loss: 0.0001
        # Step [900/938], Loss: 0.0001
        # Epoch [2/20], Train loss: 0.0008, Train Accuracy: 99.93%
        # Epoch 2 completed in 11.75 seconds
        # Test loss: 0.0001, Test Accuracy: 97.93%
        # Step [0/938], Loss: 0.0025
        # Step [100/938], Loss: 0.0032
        # Step [200/938], Loss: 0.0011
        # Step [300/938], Loss: 0.0039
        # Step [400/938], Loss: 0.0060
        # Step [500/938], Loss: 0.0081
        # Step [600/938], Loss: 0.0081
        # Step [700/938], Loss: 0.0081
        # Step [800/938], Loss: 0.0081
        # Step [900/938], Loss: 0.0081
        # Epoch [3/20], Train loss: 0.0004, Train Accuracy: 99.95%
        # Epoch 3 completed in 11.96 seconds
        # Test loss: 0.0001, Test Accuracy: 97.94%
        # Step [0/938], Loss: 0.0021
        # Step [100/938], Loss: 0.0144
        # Step [200/938], Loss: 0.0031
        # Step [300/938], Loss: 0.0081
        # Step [400/938], Loss: 0.0081
        # Step [500/938], Loss: 0.0081
        # Step [600/938], Loss: 0.0081
        # Step [700/938], Loss: 0.0081
        # Step [800/938], Loss: 0.0081
        # Step [900/938], Loss: 0.0081
        # Epoch [4/20], Train loss: 0.0009, Train Accuracy: 99.95%
        # Epoch 4 completed in 12.35 seconds
        # Test loss: 0.0001, Test Accuracy: 97.95%
        # Step [0/938], Loss: 0.0048
        # Step [100/938], Loss: 0.0056
        # Step [200/938], Loss: 0.0073
        # Step [300/938], Loss: 0.0021
        # Step [400/938], Loss: 0.0055
        # Step [500/938], Loss: 0.0055
        # Step [600/938], Loss: 0.0055
        # Step [700/938], Loss: 0.0055
        # Step [800/938], Loss: 0.0055
        # Step [900/938], Loss: 0.0055
```

Question 6

Is training faster now that it is on a GPU? Is the speedup what you would expect? Why or why not? Edit the cell below to answer.

Training on the GPU is generally expected to be faster compared to the CPU, especially when working with larger models and datasets. In this case, since we're training a relatively small MLP model on the MNIST dataset, the speedup on a GPU might not be extremely noticeable compared to training on a CPU.

Another Model Type: CNN

Until now you have trained a simple MLP for MNIST classification, however, MLPs are not a particularly good for images.

Firstly, using a MLP will require that all images have the same size and shape, since they are unrolled in the input.

Secondly, in general images can make use of translation invariance (a type of data symmetry), but this cannot be leveraged with a MLP.

For these reasons, a convolutional network is more appropriate, as it will pass kernels over the 2D image, removing the requirement for a fixed image size and leveraging the translation invariance of the 2D images.

Let's define a simple CNN below.

```
# Define the CNN model
class CNN(nn.Module):
    # Define the constructor for the network
    def __init__(self):
        super().__init__()
        # Instead of declaring the layers independently, let's use the nn.Sequential feature
```

```

# these blocks will be executed in list order

# you will break up the model into two parts:
# 1) the convolutional network
# 2) the prediction head (a small MLP)

# the convolutional network
self.net = nn.Sequential(
    nn.Conv2d(1, 32, kernel_size=3, stride=1, padding=1), # the input projection layer - note that a stride of 1 means you are not down-sampling
    nn.ReLU(), # activation
    nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1), # an inner layer - note that a stride of 2 means you are down sampling. The output is 28x28 -> 14x14
    nn.ReLU(),
    nn.Conv2d(64, 128, kernel_size=3, stride=2, padding=1), # another inner layer - note that a stride of 2 means you are down sampling. The output is 14x14 -> 7x7
    nn.ReLU(),
    nn.Linear(7*7*128, 10), # a pooling layer which will output a 3x1 vector for the prediction head
)
)

# the prediction head
self.head = nn.Sequential(
    nn.Linear(128, 64), # input projection, the output from the pool layer is a 128 element vector
    nn.ReLU(),
    nn.Linear(64, 10) # activation
)

# define the forward pass compute graph
def forward(self, x):

    # pass the input through the convolution network
    x = self.net(x)

    # reshape the output from 8x128x1x1 to 8x128
    x = x.view(x.size(0), -1)

    # pass the pooled vector into the prediction head
    x = self.head(x)

    # the output here is 8x10
    return x

# create the model
model = CNN()

# print the model and the parameter count
print(model)
param_count = sum([p.numel() for p in model.parameters()])
print("Model has (%d,%d) trainable parameters" % (param_count, param_count))

# the loss function
criterion = nn.CrossEntropyLoss()

# then you can instantiate the optimizer. You will use Stochastic Gradient Descent (SGD), and can set the learning rate to 0.1 with a
# few easier steps of 0.5
# the first input to the optimizer is the list of model parameters, which is obtained by calling .parameters() on the model object
optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.5)

# CNN
(nn):
    (net): Sequential(
        (0): Conv2d(1, 32, kernel_size=3, stride=1, dilation=1, padding=1, bias=True)
        (1): ReLU()
        (2): Conv2d(32, 64, kernel_size=3, stride=1, dilation=1, padding=1, bias=True)
        (3): ReLU()
        (4): Conv2d(64, 128, kernel_size=3, stride=2, dilation=1, padding=1, bias=True)
        (5): ReLU()
        (6): AdaptiveAvgPool2d(output_size=1)
    )
    (head): Sequential(
        (0): Linear(in_features=128, out_features=64, bias=True)
        (1): ReLU()
        (2): Linear(in_features=64, out_features=10, bias=True)
    )
Model has 101,578 trainable parameters

```

Question 7

Notice that this model now has fewer parameters than the MLP. Let's see how it trains. Using the previous code to train on the CPU with timing, edit the cell below to execute 2 epochs of training.

```

# Create a new proxy to log the loss and accuracy
train_losses = []
train_steps = []
test_losses = []
test_steps = []
test_accuracies = []

# Set global step # Start with global step 0
current_epoch = 0 # Start with epoch 0
current_step = 0 # Start with step 0
start_time = time.time() # Record the start time for the epoch

# Train on the GPU
current_step = cnp_train(current_epoch, train_losses, train_steps, current_step)

# Test on the GPU
cnp_test(test_losses, test_accuracy, test_steps, current_step)

epoch_time = time.time() - start_time # Calculate the time for this epoch
print("Epoch %d (%d steps) completed in (%.2f seconds)" % (epoch, current_step, epoch_time))

current_epoch += 1 # Move to the next epoch

# Train Epoch 0 @ [57000/50000] (95%) Loss: 0.482281: 10000 | 918/938 [01:09:00]
Testing... 10000/10000 [15/17/2018 04:48:16] Loss: 36.7215% | 918/938 [01:12:00]

Test set: Average loss: 0.4551, Accuracy: 8464/10000 (95%)

Epoch 1 completed in 75.67 seconds
Train Epoch: 1 @ [57000/50000] (95%) Loss: 0.481732: 10000 | 918/938 [01:09:00]
Testing... 10000/10000 [15/17/2018 04:48:16] Loss: 36.7215% | 918/938 [01:12:00]

Test set: Average loss: 0.4718, Accuracy: 8133/10000 (91%)
Epoch 2 completed in 76.13 seconds

```

```
# train for 2 epochs on the CPU
# Create a new array to log the loss and accuracy
cpu_train_losses = []
cpu_train_steps = []
cpu_test_losses = []
cpu_test_steps = []
cpu_test_accuracy = []

current_step = 0 # A step with global step 0
current_epoch = 0 # A epoch with epoch 0
start_time = time.time() # Record the start time for the epoch

# Train on the CPU
cpu_train_losses, cpu_train_steps, current_step = cpu_train(current_epoch, train_losses, train_steps, current_step)

# Test on the CPU
cpu_test_losses, test_accuracy, test_steps, current_step = cpu_test(cpu_train_losses, test_accuracy, test_steps, current_step)

epoch_time = time.time() - start_time # Calculate the time for this epoch
print(f'Epoch {current_epoch + 1} completed in {epoch_time:.2f} seconds')

current_epoch += 1 # Move to the next epoch

# Final Epoch: 0 [57000/00000 (0%)]: loss: 0.137401, 1000/ 938/938 [01:00:00, 00: 13.661]
Testing.... 1000/ 938 [375/375] [08:04:00, 00: 34.411/3]
Test set: Average loss: 0.1299, Accuracy: 9458/10000 (95%)
Epoch 1 completed in 73.70 seconds
# Final Epoch: 1 [57000/00000 (100%)]: loss: 0.088054, 1000/ 938/938 [01:00:00, 00: 13.771]
Testing.... 1000/ 938 [375/375] [08:04:00, 00: 33.211/3]
Test set: Average loss: 0.1299, Accuracy: 9906/10000 (96%)
Epoch 2 completed in 73.78 seconds
```

Question 8

Now, let's move the model to the GPU and try training for 2 epochs the

```

# create the model
model = CNN()

model.compile()

# print the model, and the parameter count
print(model)
for p, v in zip(model.get_params(), model.get_params()):
    print(p.name, v)

print("Model has (%d, %d) trainable parameters" % (len(model.trainable_weights), len(model.non_trainable_weights)))

# the loss Function
criterion = nn.CrossEntropyLoss()

# then you can instantiate the optimizer. You will use Stochastic Gradient Descent (SGD), and can set the learning rate to 0.1 with a momentum factor of 0.9
# the first Input to the optimizier is the list of model parameters, which is obtained by calling .parameters() on the model object
optimizer = optim.SGD(model.parameters(), lr=0.1, momentum=0.5)

# CNN
class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        # Conv2d(1, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        self.conv1 = nn.Conv2d(1, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        # ReLU()
        self.relu1 = nn.ReLU()
        # Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
        self.conv2 = nn.Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
        # ReLU()
        self.relu2 = nn.ReLU()
        # Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
        self.conv3 = nn.Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
        # ReLU()
        self.relu3 = nn.ReLU()
        # AdaptiveAvgPool2d(output_size=1)
        self.adaptive_avg_pool1 = nn.AdaptiveAvgPool2d((1, 1))
        # Linear(in_features=128, out_features=64, bias=True)
        self.linear1 = nn.Linear(128, 64)
        # ReLU()
        self.relu4 = nn.ReLU()
        # Linear(in_features=64, out_features=10, bias=True)
        self.linear2 = nn.Linear(64, 10)
        # ReLU()
        self.relu5 = nn.ReLU()

    def forward(self, x):
        x = self.conv1(x)
        x = self.relu1(x)
        x = self.conv2(x)
        x = self.relu2(x)
        x = self.conv3(x)
        x = self.relu3(x)
        x = self.adaptive_avg_pool1(x)
        x = x.view(x.size(0), -1)
        x = self.linear1(x)
        x = self.relu4(x)
        x = self.linear2(x)
        x = self.relu5(x)
        return x

Model has 181,728 trainable parameters

# create a new array to log the loss and accuracy
train_losses = []
train_accuracies = []
test_losses = []
test_accuracies = []

# Set the number of epochs
num_epochs = 10

# Set the learning rate
learning_rate = 0.001

# Start with global step 0
current_global_step = 0
current_epoch = 0
start_time = time.time()

# Start training on the GPU for 2 epochs
# train_losses, train_accuracies = train(model, criterion, optimizer, num_epochs, learning_rate, current_global_step, start_time)
# test_losses, test_accuracies = validate(model, criterion, num_epochs, learning_rate, current_global_step, start_time)
# print("Training complete! Took %d seconds." % (time.time() - start_time))

```

```

start_time = time.time() # Record the start time for the epoch
# Train on the GPU
current_step = gpu_train(current_epoch, train_losses, train_steps, current_step)
# Test on the GPU
gpu_test(test_losses, test_accuracy, test_steps, current_step)
epoch_time = time.time() - start_time # calculate the time for this epoch
print(f"Epoch {current_epoch + 1} completed in {epoch_time:.2f} seconds")
current_epoch += 1 # Move to the next epoch

```

[1] Step [0/938], loss: 2.3044

Step [100/938], loss: 2.2487

Step [200/938], loss: 2.1975

Step [300/938], loss: 2.1552

Step [400/938], loss: 2.0611

Step [500/938], loss: 1.9763

Step [600/938], loss: 1.8783

Step [700/938], loss: 1.7183

Step [800/938], loss: 1.5787

Step [900/938], loss: 1.4387

Epoch 1/20, Train Loss: 1.5115, Train Accuracy: 48.5%

Epoch 1 completed in 16.01 seconds

Test 1 completed in 16.01 seconds, 71.5%

Epoch 1 completed in 18.01 seconds

Step [0/938], loss: 0.6338

Step [100/938], loss: 0.2891

Step [200/938], loss: 0.1939

Step [300/938], loss: 0.1719

Step [400/938], loss: 0.1422

Step [500/938], loss: 0.1122

Step [600/938], loss: 0.0409

Epoch 2/20, Train Loss: 0.3612, Train Accuracy: 89.54%

Epoch 2 completed in 16.01 seconds

Test Loss: 0.2787, Test Accuracy: 91.17%

Epoch 2 completed in 16.76 seconds

train for 2 epochs on the GPU

Create a new array to log the loss and accuracy

train_losses = []

test_losses = []

test_accuracy = []

current_step = 0 # Start with global step 0

current_epoch = 0 # Start with epoch 0

A Move the model to the GPU

model.cuda()

A Start training on the GPU for 2 epochs

for epoch in range(0, 2): # Train for 2 epochs
 start_time = time.time() # Record the start time for the epoch

Train on the GPU
 current_step = gpu_train(current_epoch, train_losses, train_steps, current_step)

Test on the GPU
 gpu_test(test_losses, test_accuracy, test_steps, current_step)

epoch_time = time.time() - start_time # calculate the time for this epoch

print(f"Epoch {current_epoch + 1} completed in {epoch_time:.2f} seconds")

current_epoch += 1 # Move to the next epoch

[2] Step [0/938], loss: 0.2061

Step [100/938], loss: 0.2267

Step [200/938], loss: 0.1242

Step [300/938], loss: 0.0982

Step [400/938], loss: 0.0892

Step [500/938], loss: 0.0752

Step [600/938], loss: 0.1052

Step [700/938], loss: 0.1693

Step [800/938], loss: 0.1188

Epoch 1/20, Train Loss: 0.2064, Train Accuracy: 93.72%

Epoch 1 completed in 16.01 seconds

Test Loss: 0.1792, Test Accuracy: 94.19%

Epoch 2/20, Train Loss: 0.1515, Train Accuracy: 95.33%

Epoch 2 completed in 14.79 seconds

Test Loss: 0.1584, Test Accuracy: 94.01%

Epoch 2 completed in 16.01 seconds

Question 9

How do the CPU and GPU versions compare for the CNN? Is one faster than the other? Why do you think this is, and how does it differ from the MLP?

Edit the cell below to answer.

Double-click (or enter) to edit

As a final comparison, you can profile the FLOPs (floating-point operations) executed by each model. You will use the thop.profile function for this and consider an MNIST batch size of 1.

```

# the input shape of a MNIST sample with batch_size = 1
input = torch.randn(1, 1, 28, 28)

# create a copy of the model on the CPU
mlp_model = MLP()
cpu_model = copy.deepcopy(mlp_model)

# profile the MLP
flops, params = thop.profile(mlp_model, inputs=(input,), verbose=False)
print("MLP has {} parameters and uses {} FLOPs".format(params, flops))

# profile the CNN
flops, params = thop.profile(cpu_model, inputs=(input,), verbose=False)
print("CNN has {} parameters and uses {} FLOPs".format(params, flops))

```

Question 10

Are these results what you would have expected? Do they explain the performance difference between running on the CPU and GPU? Why or why not? Edit the cell below to answer.

The profiling results are expected. CNNs have fewer parameters than MLPs but may require more FLOPs due to convolution operations. While CPUs are general-purpose and lend a hand to tasks like convolutions, GPUs excel at parallel computations, making CNNs faster to train on GPUs. The performance difference arises because CNNs leverage parallelism in GPUs better than MLPs, which have fully connected layers and don't benefit as much from GPU parallelism.